



Step 1 - Types of Databases



What all we'll learn today –

Simple – SQL vs NoSQL, how to create Postgres Databases,

How to do CRUD on them

Advance – Relationships, Joins, Transactions

There are a few types of databases, all serve different types of use-cases

NoSQL databases

1. Store data in a **schema-less** fashion. Extremely lean and fast way to store data.
2. Examples – MongoDB,

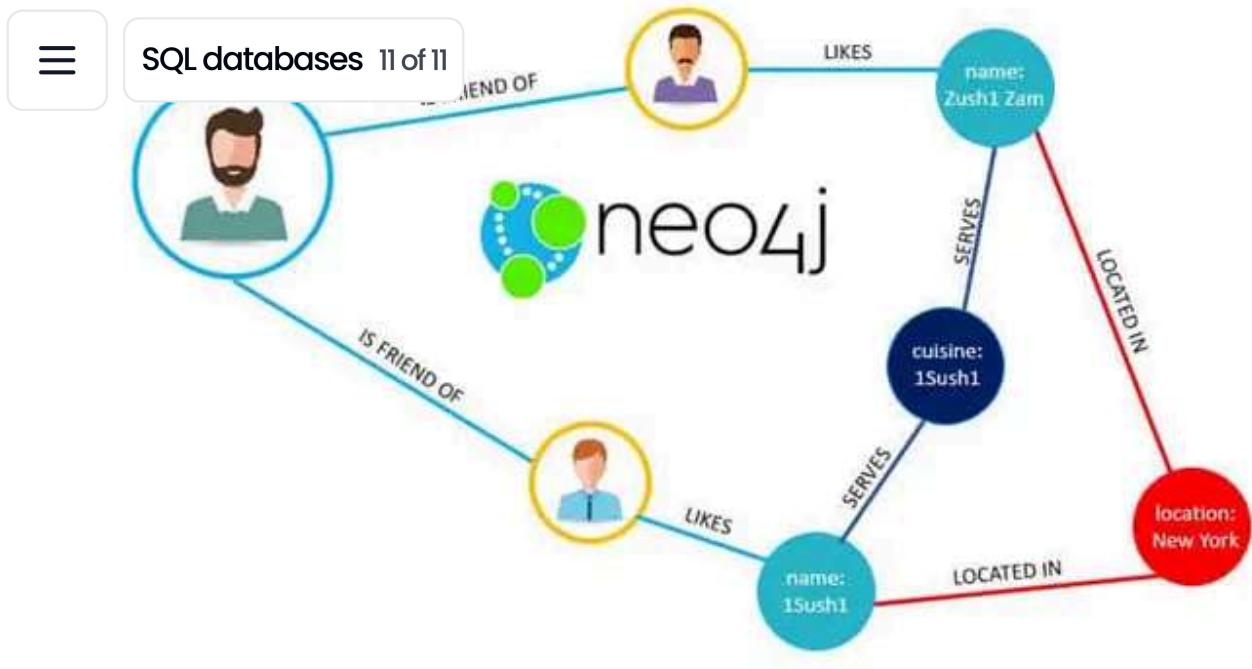


The screenshot shows the MongoDB Compass interface. At the top, there's a header with 'populations.cities' and 'Documents'. Below the header, the title 'SQL databases 11 of 11' is displayed, followed by a navigation bar with 'Documents', 'Aggregations', 'Schema', 'Explain Plan', 'Indexes', and 'Validation'. The 'Documents' tab is selected. A search bar contains the query '{ field: 'value' }'. To the right of the search bar are buttons for 'OPTIONS', 'FIND', 'RESET', and a refresh icon. Below the search bar, a message says 'Displaying documents 1 - 20 of 20' with navigation arrows. There are four document cards listed:

- `_id: ObjectId("617604944e16573d5867039b")
name: "Seoul"
country: "South Korea"
continent: "Asia"
population: 25.674`
- `_id: ObjectId("617604944e16573d5867039c")
name: "Mumbai"
country: "India"
continent: "Asia"
population: 19.98`
- `_id: ObjectId("617604944e16573d5867039d")
name: "Lagos"
country: "Nigeria"
continent: "Africa"
population: 13.463`
- `_id: ObjectId("617604944e16573d5867039e")`

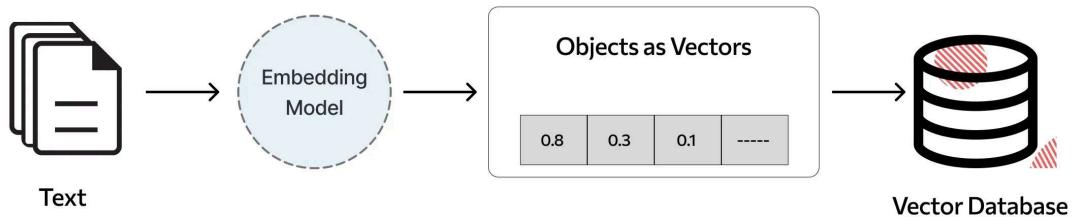
Graph databases

1. Data is stored in the form of a graph. Specially useful in cases where **relationships** need to be stored (social networks)
2. Examples – Neo4j



Vector databases

1. Stores data in the form of vectors
2. Useful in Machine learning
3. Examples – Pinecone



oraft



1. Stores data in the form of rows

Step 2 – Why not NoSQL

You might've used [MongoDB](#)

It's [schemaless](#) properties make it ideal to for bootstrapping a project fast.

But as your app grows, this property makes it very easy for data to get [corrupted](#)

What is schemaless?

Different rows can have different [schema](#) (keys/types)



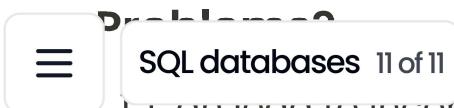
```
1  _id: ObjectId('65b3f3319e01786d275501c8')
2  userId: 65b3f3319e01786d275501c6
3  balance: 5885.875967139374
4  __v: 0
```



```
1  _id: ObjectId('65b3f3469e01786d275501ce')
2  userId: ObjectId('65b3f3469e01786d275501cc')
3  __v: 0
4  amountBalance: 720.3759487457523
```



```
1  _id: ObjectId('65b3f3499e01786d275501d3')
2  userId: ObjectId('65b3f3499e01786d275501d1')
3  balance: '"harkirat"'
```



1. Can lead to inconsistent database
2. Can cause runtime errors
3. Is too flexible for an app that needs strictness

Upsides?

1. Can move very fast
2. Can change schema very easily



You might think that `mongoose` does add strictness to the codebase because we used to define a schema there. That strictness is present at the Node.js level, not at the DB level. You can still put in erroneous data in the database that doesn't follow that schema.

Step 3 – Why SQL?

na. They require you to

2. Put in data that follows that schema



SQL databases 11 of 11

na as your app changes and perform
migrations

So there are 4 parts when using an SQL database (not connecting it to Node.js, just running it and putting data in it)

1. Running the database.
2. Using a library that lets you connect and put data in it.
3. Creating a table and defining its **schema**.
4. Run queries on the database to interact with the data
(Insert/Update/Delete)

Step 4 - Creating a database

in a few ways -

SQL databases 11 of 11

<https://neon.tech/> is a decent service that let's you create a server.

The screenshot shows the Neon.tech dashboard with a single project named 'test'. The 'Projects' table has columns for Name, Region, Created at, Storage, and Postgres version. The 'test' project is listed with US East (Ohio) as the region, created on Sep 3, 2023 at 4:23 pm, 99 MiB of storage, and Postgres version 15. Below the table, the 'Connection Details' section displays the connection string: `postgresql://100xdevs:*****@ep-broken-frost-69135494.us-east-2.ams.neon.tech/calm-gobbler-41_db_2253874?sslmode=require`. A red arrow points from the 'test' row in the table to this connection string. To the right of the connection details, there are sections for 'Branches' and 'Usage since January 1'.

Connection String

`postgresql://username:password@ep-broken-frost-69135494.us-east-2.ams.neon.tech/calm-gobbler-41_db_2253874?sslmode=require`

▼ Using docker locally

`psql -e POSTGRES_PASSWORD=mysecretpassword`

The screenshot shows a browser window with a header bar containing three horizontal lines, the text "SQL databases 11 of 11", and the word "ring". Below the header is a URL bar with the text "postgresql://postgres:mysecretpassword@localhost:5432/postgres". Underneath the URL bar is a navigation bar with left and right arrows. The main content area contains a section titled "▼ Using docker on windows" followed by a list of instructions and command examples.

▼ Using docker on windows

How to run postgreSQL in windows terminal(if you have docker installed).

- first run docker gui application that help in running commands in terminal.
- After that run it with the docker instance by the help of following command .
 - for the first time if the image is not downloaded .
 - `docker run --name my-postgres1 -e POSTGRES_PASSWORD=mysecretpassword -d -p 5432:5432 postgres` .
 - if the docker image is there, prior to use the it can simply be runned by `docker run <image name>` .
- After that ,
 - use `docker exec -it my-postgres1 psql -U postgres -d postgres` this command in terminal .
 - then enter the password and it will connect to localhost Postgress instance .
 - now you will be inside the postress command line that looks like `postgres-#` .
- U can check it by running `\dt` , (the command to display all the tables.)



The connection string is similar to the string we had in mongoose.

Connection String



`postgres://username:password@host/database`

SQL databases 11 of 11

where does harkirat live ⇒ [1, 2, 2, 2, 3001, 100]

Harkirat lives in INdia ⇒ [1, 2, 2, 2, 2, 2]

Harkirat is from chandigarh ⇒ [1, 2, 2, 2, 3]

Harkirat has been living in india, chandigarh ⇒ [1, 2, 2, 2, 2, 3]

The world is round ⇒ [1, 2, 10001, 1001, 001001]

Pacman is such a good game ⇒ [100, 10001, 20020, 1-001, 100]

Step 5 – Using a library that let's you connect and put data in it.

1. psql

`psql` is a terminal-based front-end to PostgreSQL. It provides an interactive command-line interface to the PostgreSQL (or TimescaleDB) database. With `psql`, you can type in queries interactively, issue them to PostgreSQL, and see the query results.

base?

psql Comes bundled with postgresql. You don't need it for this

SQL databases 11 of 11 tly be communicating with the database from
Node.js

```
psql -h p-broken-frost-69135494.us-east-2.aws.neon.tech -d data
```

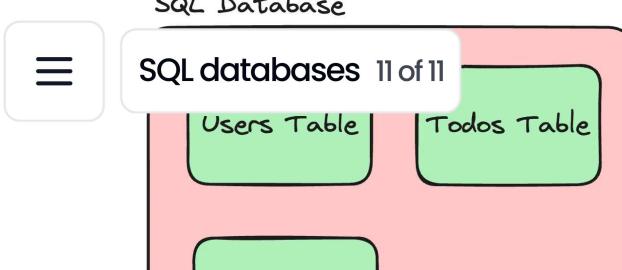
2. pg

pg is a Node.js library that you can use in your backend app to store data in the Postgres DB (similar to mongoose). We will be installing this eventually in our app.

Step 6 - Creating a table and defining its schema .

Tables in SQL

A single database can have multiple tables inside. Think of them as collections in a MongoDB database.



Until now, we have a database that we can interact with. The next step in case of postgres is to define the **schema** of your tables.

SQL stands for **Structured query language**. It is a language in which you can describe what/how you want to put data in the database.

To create a table, the command to run is -

```
CREATE TABLE users (
    id SERIAL PRIMARY KEY,
    username VARCHAR(50) UNIQUE NOT NULL,
    email VARCHAR(255) UNIQUE NOT NULL,
    password VARCHAR(255) NOT NULL,
    created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
);
```

There are a few parts of this SQL statement, let's decode them one by one

1. CREATE TABLE users

CREATE TABLE users : This command initiates the creation of a new table in the database named **users**.

2. id SERIAL PRIMARY KEY

- **id** : The name of the first column in the **users** table, typically used as a unique identifier for each row (user). Similar to **_id** in mongodb

data type for creating an auto-incrementing integer. Every time a new row is inserted, this

value automatically increments, ensuring each user has a



SQL databases 11 of 11

- **PRIMARY KEY** : This constraint specifies that the `id` column is the primary key for the table, meaning it uniquely identifies each row. Values in this column must be unique and not null.

3. email VARCHAR(255) UNIQUE NOT NULL,

- `email` : The name of the second column, intended to store the user's username.
- **VARCHAR(50)** : A variable character string data type that can store up to 50 characters. It's used here to limit the length of the username.
- **UNIQUE** : This constraint ensures that all values in the `username` column are unique across the table. No two users can have the same username.
- **NOT NULL** : This constraint prevents null values from being inserted into the `username` column. Every row must have a username value.

4. password VARCHAR(255) NOT NULL

Same as above, can be non unique

5. created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP

- `created_at` : The name of the fifth column, intended to store the timestamp when the user was created.
- **TIMESTAMP WITH TIME ZONE** : This data type stores both a timestamp and a time zone, allowing for the precise tracking of timestamps regardless of the user's or server's time

- **DEFAULT CURRENT_TIMESTAMP** : This default value automatically SQL databases 11 of 11 at column to the date and time at which the row is inserted into the table, using the current timestamp of the database server.



If you have access to a database right now, try running this command to create a simple table in there

```
CREATE TABLE users (
    id SERIAL PRIMARY KEY,
    username VARCHAR(50) UNIQUE NOT NULL,
    email VARCHAR(255) UNIQUE NOT NULL,
    password VARCHAR(255) NOT NULL,
    created_at TIMESTAMP WITH TIME ZONE DEFAULT
    CURRENT_TIMESTAMP
);
```

Then try running

```
\dt;
```

to see if the table has been created or not



Step 7 - Interacting with the database

There are 4 things you'd like to do with a database

1. INSERT

```
INSERT INTO users (username, email, password)  
VALUES ('username_here', 'user@example.com', 'user_password');
```



Notice how you didn't have to specify the `id` because it auto increments

2. UPDATE

```
UPDATE users  
SET password = 'new_password'  
WHERE email = 'user@example.com';
```



3. DELETE

```
DELETE FROM users  
WHERE id = 1;
```



The screenshot shows a PostgreSQL database interface. At the top, there's a bar with a menu icon, the text "SELECT * FROM users", and a search bar containing "SQL databases 11 of 11". Below this is a table with 11 rows of data. A yellow lightbulb icon with a tooltip provides instructions on how to run the query in a terminal.

 Try running all 4 of these in your terminal if you have `psql` installed locally.
If not, that's fine we'll eventually be doing these through the `pg` library

Step 8 – How to do queries from a Node.js app?

In the end, postgres exposes a protocol that someone needs to talk to be able to send these commands (update, delete) to the database.

`psql` is one such library that takes commands from your terminal and sends it over to the database.

To do the same in a Node.js , you can use one of many `Postgres clients`

`pg` library

Non-blocking PostgreSQL client for Node.js.



SQL databases 11 of 11 <https://node-postgres.com/>

Connecting -

```
import { Client } from 'pg'

const client = new Client({
  host: 'my.database-server.com',
  port: 5334,
  database: 'database-name',
  user: 'database-user',
  password: 'secretpassword!!',
})

client.connect()
```



Querying -

```
const result = await client.query('SELECT * FROM USERS;')
console.log(result)
```



// write a function to create a users table in your database.

```
import { Client } from 'pg'
```



```
const client = new Client({
  connectionString: "postgresql://postgres:mysecretpassword@local"
})
```

```
async function createUsersTable() {
```

```
  await client.connect()
```

```
  const result = await client.query(`
```

```
    CREATE TABLE users (
```

email VARCHAR(255) UNIQUE NOT NULL,

name VARCHAR(255) UNIQUE NOT NULL,

```
password VARCHAR(255) NOT NULL,
```



SQL databases 11 of 11

```
TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIM
```

```
)  
    console.log(result)  
}
```

```
createUsersTable();
```

Step 9 – Creating a simple Node.js app

1. Initialise an empty typescript project

```
npm init -y  
npx tsc --init
```



1. Change the `rootDir` and `outDir` in `tsconfig.json`

```
"rootDir": "./src",  
"outDir": "./dist",
```



uses (because we're using TS)



npm install pg
SQL databases 11 of 11 s/pg



1. Create a simple Node.js app that lets you put data

Create a function that let's you insert data into a table. Make it `async`, make sure `client.connect` resolves before u do the insert

▼ Answer

```
import { Client } from 'pg';

// Async function to insert data into a table
async function insertData() {
  const client = new Client({
    host: 'localhost',
    port: 5432,
    database: 'postgres',
    user: 'postgres',
    password: 'mysecretpassword',
  });

  try {
    await client.connect(); // Ensure client connection is established
    const insertQuery = "INSERT INTO users (username, email, passw
    const res = await client.query(insertQuery);
    console.log('Insertion success:', res); // Output insertion result
  } catch (err) {
    console.error('Error during the insertion:', err);
  } finally {
    await client.end(); // Close the client connection
  }
}

insertData();
```



This is an insecure way to store data in your tables.

...nality eventually via `HTTP`,



someone can do an **SQL INJECTION** to get access to your **SQL databases** 11 of 11 ur data.

1. More secure way to store data.

Update the code so you don't put **user provided fields** in the **SQL string**

▼ What are user provided strings?

In your final app, this insert statement will be done when a user signs up on your app.

Email, username, password are all user provided strings

▼ What is the **SQL string** ?

```
const insertQuery = "INSERT INTO users (username, email, passw□
```

`◀ ▶`

▼ Hint

```
const insertQuery = 'INSERT INTO example_table(column1, colum□  
const res = await client.query(insertQuery, [column1Value, colum□
```

`◀ ▶`

▼ Solution

```
import { Client } from 'pg';□
```

```
// Async function to insert data into a table  
async function insertData(username: string, email: string, passwc  
const client = new Client({  
  host: 'localhost',  
  port: 5432,  
  database: 'postgres',  
  user: 'postgres',  
  password: 'mysecretpassword',  
});
```

ure client connection is establishe

```
// Use parameterized query to prevent SQL injection
SQL databases 11 of 11
query = "INSERT INTO users (username, email, password) VALUES ($1, $2, $3);"
const values = [username, email, password];
const res = await client.query(query, values);
console.log('Insertion success:', res); // Output insertion result
} catch (err) {
  console.error('Error during the insertion:', err);
} finally {
  await client.end(); // Close the client connection
}
}

// Example usage
insertData('username5', 'user5@example.com', 'user_password').then(() =>
```

1. Query data

Write a function `getUser` that lets you fetch data from the database given a `email` as input.

```
import { Client } from 'pg';
```



```
// Async function to fetch user data from the database given an email
async function getUser(email: string) {
  const client = new Client({
    host: 'localhost',
    port: 5432,
    database: 'postgres',
    user: 'postgres',
    password: 'mysecretpassword',
  });
}
```

```
try {
  await client.connect(); // Ensure client connection is established
  const query = 'SELECT * FROM users WHERE email = $1';
  const values = [email];
  const result = await client.query(query, values);
}
```

https://projects.100xdevs.com/pdf/YOSAherHkqWXhOdIE4yE/sql-11

```
console.log('User found:', result.rows[0]); // Output user data
ws[0]; // Return the user data
SQL databases 11 of 11

    console.log('No user found with the given email.');
    return null; // Return null if no user was found
}
} catch (err) {
    console.error('Error during fetching user:', err);
    throw err; // Rethrow or handle error appropriately
} finally {
    await client.end(); // Close the client connection
}
}

// Example usage
getUser('user5@example.com').catch(console.error);
```

Step 10 - Relationships and Transactions

Relationships let you store data in different tables and **relate** it with each other.

odb

Since `mongodb` is a NoSQL database, you can store any shape of

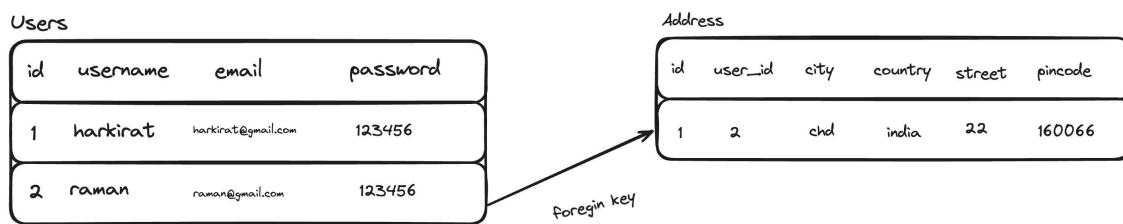
`SQL databases` 11 of 11

If I ask you to store a user's details along with their address, you can store it in an object that has the address details.

```
_id: ObjectId('65b3f3469e01786d275501ce')
address: Object
  street: "1 West HW"
  city: "Chandigarh"
  country: "India"
  pincode: "160066"
email: "harkirat@gmail.com"
name: "harkirat"
```

Relationships in SQL

Since SQL can not store `objects` as such, we need to define two different tables to store this data in.



This is called a `relationship`, which means that the `Address` table is

When defining the tables, you need to define the `relationship`

```
CREATE TABLE users (
    SQL databases 11 of 11 PRIMARY KEY,
    username VARCHAR(50) UNIQUE NOT NULL,
    email VARCHAR(255) UNIQUE NOT NULL,
    password VARCHAR(255) NOT NULL,
    created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
);

CREATE TABLE addresses (
    id SERIAL PRIMARY KEY,
    user_id INTEGER NOT NULL,
    city VARCHAR(100) NOT NULL,
    country VARCHAR(100) NOT NULL,
    street VARCHAR(255) NOT NULL,
    pincode VARCHAR(20),
    created_at TIMESTAMP WITH TIME ZONE DEFAULT CURRENT_TIMESTAMP
    FOREIGN KEY (user_id) REFERENCES users(id) ON DELETE CASCADE
);
```



SQL query

To insert the address of a user –

```
INSERT INTO addresses (user_id, city, country, street, pincode)
VALUES (1, 'New York', 'USA', '123 Broadway St', '10001');
```



Now if you want to get the address of a user given an `id` , you can run the following query –

```
SELECT city, country, street, pincode
FROM addresses
WHERE user_id = 1;
```



Extra – Transactions in SQL



SQL databases 11 of 11

to have at this point is what queries are run signs up and sends both their information and their address in a single request.

Do we send two SQL queries into the database? What if one of the queries (address query for example) fails?

This would require **transactions** in SQL to ensure either both the user information and address goes in, or neither does

▼ SQL Query

```
BEGIN; -- Start transaction
```



```
INSERT INTO users (username, email, password)  
VALUES ('john_doe', 'john_doe@example.com', 'securepassword1')
```

```
INSERT INTO addresses (user_id, city, country, street, pincode)  
VALUES (currval('users_id_seq'), 'New York', 'USA', '123 Broadway St')
```

```
COMMIT;
```



▼ Node.js Code

```
import { Client } from 'pg';
```



```
async function insertUserAndAddress(
```

```
    username: string,  
    email: string,  
    password: string,  
    city: string,  
    country: string,  
    street: string,  
    pincode: string
```

```
) {
```

```
    const client = new Client({  
        host: 'localhost',  
        port: 5432,
```

```
        user: 'postgres',
```

```
password: 'mysecretpassword',
```



SQL databases 11 of 11

```
try {
    await client.connect();

    // Start transaction
    await client.query('BEGIN');

    // Insert user
    const insertUserText = `
        INSERT INTO users (username, email, password)
        VALUES ($1, $2, $3)
        RETURNING id;
    `;
    const userRes = await client.query(insertUserText, [username, password]);
    const userId = userRes.rows[0].id;

    // Insert address using the returned user ID
    const insertAddressText = `
        INSERT INTO addresses (user_id, city, country, street, pincode)
        VALUES ($1, $2, $3, $4, $5);
    `;
    await client.query(insertAddressText, [userId, city, country, street, pincode]);

    // Commit transaction
    await client.query('COMMIT');

    console.log('User and address inserted successfully');
} catch (err) {
    await client.query('ROLLBACK'); // Roll back the transaction or
    console.error('Error during transaction, rolled back.', err);
    throw err;
} finally {
    await client.end(); // Close the client connection
}
```



'johndoe',
SQL databases 11 of 11 ample.com',
'ord123',
'New York',
'USA',
'123 Broadway St',
'10001'
);

Step 11 – Joins

Defining relationships is easy.

What's hard is **joining** data from two (or more) tables together.

For example, if I ask you to fetch me a user's details **and** their address, what SQL would you run?

▼ Approach 1 (Bad)

-- Query 1: Fetch user's details



```
SELECT id, username, email  
FROM users  
WHERE id = YOUR_USER_ID;
```

-- Query 2: Fetch user's address

```
SELECT city, country, street, pincode  
FROM addresses  
WHERE user_id = YOUR_USER_ID;
```



```
SELECT users.id, users.username, users.email, addresses.city, addresses.country
FROM users
JOIN addresses ON users.id = addresses.user_id
WHERE users.id = '1';

SELECT u.id, u.username, u.email, a.city, a.country, a.street, a.pincode
FROM users u
JOIN addresses a ON u.id = a.user_id
WHERE u.id = YOUR_USER_ID;
```

Now try converting the same to your node app

▼ Approach 1 (Bad)

```
import { Client } from 'pg';

// Async function to fetch user details and address separately
async function getUserDetailsAndAddressSeparately(userId: string) {
  const client = new Client({
    host: 'localhost',
    port: 5432,
    database: 'postgres',
    user: 'postgres',
    password: 'mysecretpassword',
  });

  try {
    await client.connect();

    // Fetch user details
    const userDetailsQuery = 'SELECT id, username, email FROM users';
    const userDetails = await client.query(userDetailsQuery, [userId]);

    // Fetch user address
    const userAddressQuery = 'SELECT city, country, street, pincode
    client.query(userAddressQuery, [userId]);
```



SQL databases 11 of 11

```

if (userDetails.rows.length > 0) {
    log('User found:', userDetails.rows[0]);
    log('Address:', userAddress.rows.length > 0 ? userAddress.rows[0].address : null);
    return { user: userDetails.rows[0], address: userAddress.rows[0] };
} else {
    console.log('No user found with the given ID.');
    return null;
}
} catch (err) {
    console.error('Error during fetching user and address:', err);
    throw err;
} finally {
    await client.end();
}
}

getUserDetailsAndAddressSeparately("1");

```

▼ Approach 2 (using joins)

```
import { Client } from 'pg';
```



```

// Async function to fetch user data and their address together
async function getUserDetailsWithAddress(userId: string) {
    const client = new Client({
        host: 'localhost',
        port: 5432,
        database: 'postgres',
        user: 'postgres',
        password: 'mysecretpassword',
    });

    try {
        await client.connect();
        const query = `
            SELECT u.id, u.username, u.email, a.city, a.country, a.street, a.zip_code
            FROM users u
            JOIN addresses a ON u.id = a.user_id
        `;

        const result = await client.query(query, [userId]);
    }
}
```



SQL databases 11 of 11

```
    vs.length > 0) {
      log('User and address found:', result.rows[0]);
      return result.rows[0];
    } else {
      console.log('No user or address found with the given ID.');
      return null;
    }
  } catch (err) {
  console.error('Error during fetching user and address:', err);
  throw err;
} finally {
  await client.end();
}
}

getUserDetailsWithAddress("1");
```

No

Benefits of using a join -

1. Reduced Latency
2. Simplified Application Logic
3. Transactional Integrity

Types of Joins

1. INNER JOIN

Returns rows when there is at least one match in both tables. If

it returned. It's the most common

Use Case: Find All Users With Their Addresses. If a user hasn't

☰ **i SQL databases 11 of 11** that user shouldn't be returned

```
SELECT users.username, addresses.city, addresses.country, address
FROM users
INNER JOIN addresses ON users.id = addresses.user_id;
```

2. LEFT JOIN

Returns all rows from the left table, and the matched rows from the right table.

Use case - To list all users from your database along with their address information (if they've provided it), you'd use a LEFT JOIN. Users without an address will still appear in your query result, but the address fields will be NULL for them.

```
SELECT users.username, addresses.city, addresses.country, address
FROM users
LEFT JOIN addresses ON users.id = addresses.user_id;
```

3. RIGHT JOIN

Returns all rows from the right table, and the matched rows from the left table.

Use case - Given the structure of the database, a RIGHT JOIN would be less common since the **addresses** table is unlikely to have entries not linked to a user due to the foreign key constraint. However, if you had a situation where you start with the **addresses** table and optionally include user information, this would be the theoretical use case.

```
SELECT users.username, addresses.city, addresses.country, address
FROM users
RIGHT JOIN addresses ON users.id = addresses.user_id;
```

4. FULL JOIN

SQL databases 11 of 11 Here is a match in one of the tables. It effectively combines the results of both LEFT JOIN and RIGHT JOIN.

Use case - A FULL JOIN would combine all records from both **users** and **addresses**, showing the relationship where it exists. Given the constraints, this might not be as relevant because every address should be linked to a user, but if there were somehow orphaned records on either side, this query would reveal them.

```
SELECT users.username, addresses.city, addresses.country, address
FROM users
FULL JOIN addresses ON users.id = addresses.user_id;
```

